

This document proposes a initial reference implementation of **archupgrade**, a software package that aims to provide reliable, predictable versioning upgrades for Arch Linux-based distributions

# 1. Overview

---

## archupgrade system requirements

For system builders:

- A build system using **archupgrade** to perform image-based, read-only system image preparation for use in deployment **should not** include **archupgrade** in the final image. Upgrades **should** be handled by the builder's choice of version control/imaging solution.

For client systems:

- A system that will be using **archupgrade** to mandate system updates **shall not** have any user-accessible means of installing packages into system repositories by default. This is, by design, so that the system states remain deterministic before and after an update is executed.
- Any such utilities that the user expects **shall** be symlinked towards the **archupgrade** binary. This is so **archupgrade** can inform the user appropriately when there is an attempt to execute such an utility.

## archupgrade specification requirements

- A specification for **archupgrade** **shall** include every single packages that will be installed on a given system. Do not assume any set of base packages will be available or will be installed as dependencies
- A specification for **archupgrade** **should** contain multiple phases, when applicable. After installation of any given phase, the system being services **should** remain fully-functional, minimizing the possibility of an unexpected error during the following phases causing complete system failures.

## archupgrade implementations specifics

- An implementation of **archupgrade** **should** be able to be executed without using any system-dependant libraries, ie. the implementation's binary should have all of its dependencies statically linked.
- An implementation of **archupgrade** **shall** be able to parse every part of a given specification, unless said part is marked as "optional"

# 2. Specification file format

---

We propose using a machine-parsable format like **json** or **yaml** for the specification file format, which shall specifies as follows:

## Versioning

A specification **shall** declare a version of the specification file format. This is required in order to enforce a minimum version for any given implementation that shall be used in order to apply said specification successfully.

The versioning **should** follow [Semantic Versioning](#). For example:

```
version: '0.0.1-alpha'
```

## Nesting specifications

A specification may choose to include other specifications as part of itself. This mechanism allows more modularity and reusability between different shipping configurations.

This is implemented as an array

```
includes:  
  - path/to/specification.yml
```

## Specifications specifics

- A specification **may** nests any arbitrary amount of other specifications. Each nested specifications may also themselves nests more specifications.
- A nested specification **shall** always be executed before its host.
- If multiple sub-specifications nests the same specification file, ie.

```
# main.yml  
includes:  
  - specification_a.yml  
  - specification_b.yml  
# specification_a.yml  
includes:  
  - specification_c.yml  
# specification_b.yml  
includes:  
  - specification_c.yml
```

said nested specifications **shall** only get executed once at the point where they are first imported (in this example, it gets executed as part of `specification_a`)

## Implementation specifics

- An implementation **shall** execute every nested specifications in the order that they are imported in

- An implementation **shall** attempt to detect and abort if nesting loops are detected in nested specifications.

## Package installation and configurations

An upgrade as specified by `archupgrade` shall run in multiple phases to account for any inconsistencies that may arise from upgrading from older package revisions. For example, it should handle anything on the [Arch Linux News](#) that would normally requires manual intervention (for example, [JDK 21](#)).

While not required, it is recommended that that each phase

For this purpose, we propose the following format:

```
upgrade:
  phases:
    first:
      backend: libalpm
      message: "Upgrading system libraries..."
      preinstall:
        - some_bash_commands
        - do_some_patching
      packages:
        - package_a:
            url:
              'https://archive.archlinux.org/p/package_a.tar.zst'
            hash: 'dQw4w9WgXcQ'
            hash-algorithm: 'sha256'
        - package_b:
            url:
              'https://archive.archlinux.org/p/package_b.tar.zst'
            hash: '_xc7tNbjnHM'
            hash-algorithm: 'sha256'
      postinstall:
        - some_more_bash_commands
        - do_some_configurations
      reboot: false
    second:
      backend: libalpm
      message: "Upgrading desktop environment..."
      preinstall:
        - some_bash_commands
        - do_some_patching
      packages:
        - package_c:
            url:
              'https://archive.archlinux.org/p/package_c.tar.zst'
            hash: 'dQw4w9WgXcQ'
            hash-algorithm: 'sha256'
        - package_d:
            url:
              'https://archive.archlinux.org/p/package_d.tar.zst'
            hash: '_xc7tNbjnHM'
```

```
        hash-algorithm: 'sha256'  
postinstall:  
    - some_more_bash_commands  
    - do_some_configurations  
reboot: true
```

## Directives

- **required-space**: Specifies an amount of space that would be necessary on the **root** directory of the target system for the upgrade to be performed successfully
- **phases**: Specifies an array of phases that the update shall be performed in, each of which includes:
  - **backend**: Specifies the backend that would be used to perform this phase of the upgrade`
  - **message**: A message that shall be displayed to the user during the phase of the upgrade
  - **preinstall**: An array of shell command that should be run before packages in this phase are installed
  - **packages**: An array of packages that shall be installed on the system. Each **package** is an object that specifies
    - **url**: A URI to the package archive itself. This may be **https://** or **file://** for remote and local archives, respectively.
    - **hash**: A hash of the package archive.
    - **hash-algorithm**: The algorithm used to generate the package's hash
  - **postinstall**: An array of shell command that should be run after packages in this phase are installed
  - **reboot**: A boolean specifying whether or not a system reboot is required before continuing with the next phase.

## Specifications specifics

- A single specification **shall not** have phases with completely matching names.
- Packages declared in any given phase **shall** provide a hash in order to protect the integrity of the downloaded package.

## Implementation specifics

- An implementation **shall** perform installation phases in order that they are specified. Packages within the same phases **may** be installed together.
- An implementation **shall** execute the **preinstall** and **postinstall** commands in their specific order before and after package installation, respectively.

- An implementation **shall** display the message as specified by `message` to the user to indicate the upgrade progress.
- An implementation **shall** perform hash calculation of the packages after obtaining them from remote or local sources.
- An implementation **may** attempt to re-download a given package archive which failed verification for a fixed number of retries. If the amount of retries exceed said fixed amounts, the implementation **shall** perform deletion of the archive to prevent accidental reuse.
- An implementation **shall** save its progress in persistent storage, in case of a temporary failure causing the upgrade process to be restarted.
- An implementation **should**, when possible, use the installed/total package counts along with the finished/total upgrade phases to provide the user with an approximate progress of installation
- An implementation **may**, when possible, attempt to perform estimations of the user system's state (eg. required disk space) that may affect the ability to complete an upgrade.

## Finalization

After an upgrade is performed by an `archupgrade` implementation, an optional finalization procedure may be performed to facilitates any post-upgrade configurations, as well as housekeeping items like cleaning up package caches or deprecated package removal.

```
finalize:
  shell:
    - some_shell_command
  file_write:
    - path: /etc/os-release
      content: >
        VERSION="10.0 (Firefly)"
        VERSION_ID=10.0
        VERSION_CODENAME=firefly
        BUILD_ID=20240618
        IMAGE_ID=firefly-shipping-final
  file_remove:
    - /etc/pacman.d/mirrorlist
  clean-caches: true
  reboot: true
```

## Directives

- `shell`: Shell commands that shall be executed after installation has completed
- `file_write`: Files that should be written to the target filesystem after installation
- `file_remove`: Files that should be removed from the target filesystem after installation

- **clean-caches**: Whether or not the implementation should remove the package cache after installation
- **reboot**: Whether a system reboot shall be issued after installation.

### Implementation specifics

- An implementation **shall** run the **shell** in the order that they appears in.
- An implementation **shall** perform writing and deletion of files in the order that they appears in. If a file already exist, the implementation **shall** overwrite said file.
- An implementation **shall** clean the package cache of every backend if they exist and has been in use when **clean-caches** is specified.